

AgentGuard Scripts

Two pieces here, for two different videos: the founder pitch (new, on camera) and the demo narration (voiceover over the video already sent).

Founder pitch video

Target length is 2.5 minutes. Deliver this on camera in your own words, not read word for word. Every factual claim below matches something real in the repo. Exact commands and captured output are in `demo/README.md` if you want to check anything.

Who I am.

[Fill in your own background here, one or two sentences: what you did before this, why this problem, why you're the one building it.]

The problem. AI agents are starting to hold wallets and sign their own Solana transactions, vendor payments, x402 agentic payments, multisig approvals. Some custodial wallets now add their own IDL-based checks, but each one locks you into that provider's engine and its own gaps, so whether a hidden instruction gets caught depends on which wallet is underneath. Not theoretical: a USENIX Security 2026 paper tested 15 production x402 facilitators and found policy violations in every single one.

What AgentGuard does. It sits in front of any agent wallet, not locked into one provider. Before a transaction signs, it decodes the registered instructions against an authenticated IDL bound to the real onchain program, so a spoofed IDL can't make a malicious instruction look benign. It checks that decoded intent against policy and either lets the transaction through automatically or holds it and routes it to a structured human review, instead of just blocking it and moving on. Anything unknown or unverified fails closed to review, never guessed at. It already gates Kora, the Solana Foundation's own relay, and reads live Squads multisig proposals to flag pending program upgrades.

Why Solana. Solana carries roughly 70% of monthly x402 transaction volume, so agents are already transacting autonomously here at real scale, which is exactly where this gap matters most.

Who the first users are. Teams shipping autonomous Solana agent wallets who need a policy gate in front of signing, and Squads multisig signers who currently get no automated context before approving a program upgrade.

[Name specific teams or communities you're already talking to, if there are any.]

What's real today. This isn't a mockup. 70 automated tests, a working decode and policy pipeline, a live read of real onchain Squads proposals. And beyond tests: a real transaction went through the full pipeline, decoded, policy checked, signed, and submitted by a live Kora relay, and it landed and finalized on Solana mainnet. Real money, a real signature anyone can look up and verify. That's proven, not just should work.

[Cut to the demo video here.]

What's next. Honestly, nobody outside this repo has used AgentGuard on a real agent wallet yet. The next milestone is getting it in front of real agent-wallet builders, running their actual traffic through it, and finding out where the policy model breaks, rather than adding more features in isolation.

[Replace this with your actual plan and timeline if it's different.]

Demo video narration

This is the narration for `agentguard-demo-video.mp4`, the video already sent to you. It's silent, 54 seconds, and built from the real output of running AgentGuard's four demo scenarios. Colosseum's cap is 3 minutes, so there's room to breathe.

How to add narration. Record yourself saying the lines below in your own words rather than reading them exactly, then lay that audio under the existing video in whatever editor is easiest for you (QuickTime, iMovie, CapCut all work). The video already has pauses timed in at each beat, so no re-cutting is needed, just talk over each section as it plays.

If you'd rather record a completely fresh version yourself instead of narrating over this one, there's a script for that too: `./demo/record-demo-video.sh` paces the real commands live in your terminal so you can screen record and narrate in a single take. That's entirely optional. The video you already have doesn't need it.

Beats

1. **Intro.** AgentGuard sits between an agent and its wallet. Before anything signs, it decodes registered instructions, checks them against policy, and routes unknown or policy-violating actions to a human.
2. **Legitimate payment.** A normal vendor payment. One transfer, decodes cleanly, matches policy, goes through.
3. **Adversarial payment.** Same payment, but a second instruction has been quietly appended that reassigns the agent's own account to an attacker-controlled program. It moves zero balance, so anything that only checks balances would miss it completely.
4. **Why it's caught.** That second instruction isn't in AgentGuard's registry, so it refuses to guess what it does and flags the whole transaction for a human instead of letting it through or blocking it blindly.
5. **Unapproved destination.** This one decodes perfectly cleanly, nothing hidden, it's just paying an address that was never approved. The kind of mistake a prompt-injected agent could make without any single instruction looking wrong.
6. **Over the limit.** Same idea, different rule. A real, approved vendor, but the amount is over the daily limit.
7. **Close.** Every one of those gets turned into a real review task for a human instead of a silent guess or a black-box block. That's AgentGuard.